

```
import pandas as pd

# Load the CSV file
file_path = "baseball_ratios.csv"
df = pd.read_csv(file_path)

#take out the categorical data
df = df.drop(columns=['year', 'team_name'])
df = df.dropna()
print(df.head())

from sklearn.preprocessing import StandardScaler

# Initialize the standard scaler
scaler = StandardScaler()

# Normalize the data
normalized_data = scaler.fit_transform(df)

normalized_df = pd.DataFrame(normalized_data, columns=df.columns)

# Display the first few rows and summary statistics of the normalized data
head_normalized = normalized_df.head()
summary_normalized = normalized_df.describe()

head_normalized, summary_normalized

from sklearn.cluster import KMeans
import matplotlib.pyplot as plt

# Define a range of clusters
cluster_range = range(1, 11)

# Compute KMeans for each number of clusters
inertias = []
for k in cluster_range:
    kmeans = KMeans(n_clusters=k, random_state=42).fit(normalized_data)
    inertias.append(kmeans.inertia_)

# Plot the Elbow method
plt.figure(figsize=(10,6))
plt.plot(cluster_range, inertias, marker='o', linestyle='--')
plt.title('Elbow Method')
plt.xlabel('Number of Clusters')
plt.ylabel('Inertia')
plt.grid(True)
```

```


plt.show()

#print inertias
for k, inertia in zip(cluster_range, inertias):
    print(f"Number of clusters: {k}, Inertia: {inertia:.2f}")

```

	rank	batting_avg	games_played	home_game_ratio	runs_scored_pg	...	\
0	5	0.293746	134	0.492537	5.671642		
1	2	0.278052	138	0.500000	5.500000		
2	3	0.286739	137	0.496350	5.430657		
3	5	0.248630	140	0.500000	3.792857		
4	1	0.275767	137	0.518248	5.978102		
	at_bats_pg	hits_pg	doubles_pg	triples_pg	homeruns_pg	...	\
0	34.246269	10.059701	1.335821	0.828358	0.179104	...	
1	35.260870	9.804348	1.326087	0.753623	0.268116	...	
2	35.613139	10.211679	1.503650	0.678832	0.233577	...	
3	33.900000	8.428571	0.964286	0.257143	0.200000	...	
4	34.489051	9.510949	1.262774	0.649635	0.233577	...	
	shutouts_pg	saves_pg	outs_pitches_pg	hits_allowed_pg	...	\	
0	0.029851	0.022388	25.925373	9.798507			
1	0.050725	0.007246	26.456522	8.536232			
2	0.051095	0.036496	26.576642	9.080292			
3	0.078571	0.000000	27.064286	8.542857			
4	0.080292	0.014599	26.678832	9.124088			
	homeruns_allowed_pg	walks_allowed_pg	strikeouts_by_pitchers_pg	...	\		
0	0.156716	2.567164		2.022388			
1	0.239130	2.130435		2.869565			
2	0.131387	3.175182		4.255474			
3	0.207143	2.492857		3.985714			
4	0.197080	2.277372		2.875912			
	errors_pg	double_plays_pg	fielding_percentage_pg				
0	2.992537	0.567164	0.006910				
1	2.442029	0.753623	0.006833				
2	2.051095	0.722628	0.006934				
3	2.014286	0.635714	0.006800				
4	2.518248	0.729927	0.006869				

[ 5 rows x 27 columns ]



```
Number of clusters: 1, Inertia: 42498.00
Number of clusters: 2, Inertia: 31697.26
Number of clusters: 3, Inertia: 27396.55
Number of clusters: 4, Inertia: 25372.85
Number of clusters: 5, Inertia: 23725.78
Number of clusters: 6, Inertia: 20811.97
Number of clusters: 7, Inertia: 19879.10
Number of clusters: 8, Inertia: 18883.56
Number of clusters: 9, Inertia: 18327.78
Number of clusters: 10, Inertia: 17688.58
```

```
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt

X = df.values

# 1. Apply PCA to reduce to 2 components
pca = PCA(n_components=2)
```

```


X_pca = pca.fit_transform(X)

# 2. Perform KMeans k=3 from before
kmeans = KMeans(n_clusters=3, random_state=42)
kmeans.fit(X_pca)
y_kmeans = kmeans.predict(X_pca)

# 3. Plot the clustered points
plt.scatter(X_pca[:, 0], X_pca[:, 1], c=y_kmeans, cmap='viridis', s=50)
centers = kmeans.cluster_centers_
plt.scatter(centers[:, 0], centers[:, 1], c='red', s=200, alpha=0.75,
plt.title('KMeans Clusters after PCA Reduction')
plt.xlabel('PCA Component 1')
plt.ylabel('PCA Component 2')
plt.show()

print(pca.explained_variance_ratio_)

```



[ 0.96331682 0.0167524 ]

```

import numpy as np
import seaborn as sns
from scipy import stats

```

```

# 1. Add cluster labels to your original dataframe
df['Cluster'] = kmeans.labels_

# 2. summary statistics for each cluster
unique_clusters = df['Cluster'].unique()

# Create a summary df with means for each feature by cluster
cluster_means = df.groupby('Cluster').mean()
cluster_medioms = df.groupby('Cluster').median()
cluster_stds = df.groupby('Cluster').std()
cluster_counts = df.groupby('Cluster').size()

print("Cluster sizes:")
print(cluster_counts)
print("\nCluster means:")
print(cluster_means)

# 3. Find top 10 features with biggest differences between clusters
# Calc the var of means across clusters for each feature
feature_variance = cluster_means.var().sort_values(ascending=False)
most_distinctive_features = feature_variance.index[:10]

print("\nTop features (highest variance between clusters):")
print(feature_variance.head(10))

# Run ANOVA to find statistically significant differences
from scipy.stats import f_oneway

anova_results = {}
feature_columns = [col for col in df.columns if col != 'Cluster']

for feature in feature_columns:
    feature_by_cluster = [df[df['Cluster'] == c][feature].dropna() for c in unique_clusters]
    # Only run ANOVA if we have sufficient data in each group
    if all(len(group) > 1 for group in feature_by_cluster):
        anova_result = f_oneway(*feature_by_cluster)
        anova_results[feature] = {
            'F-statistic': anova_result.statistic,
            'p-value': anova_result.pvalue
        }

# Convert to df and sort by F-statistic
anova_df = pd.DataFrame(anova_results).T
anova_df = anova_df.sort_values('F-statistic', ascending=False)
print("\nANOVA Results (most significant differences):")
print(anova_df.head(10))

```

dtype: int64

Cluster means:

Cluster	rank	batting_avg	games_played	home_game_ratio	runs_s	.
0	3.389986	0.258440	161.187588		0.499947	.
1	2.933333	0.244196	59.866667		0.499923	.
2	3.682540	0.264770	127.920635		0.500033	.
Cluster	at_bats_pg	hits_pg	doubles_pg	triples_pg	homeruns_pg	.
0	34.018116	8.797347	1.596714	0.220030	0.878615	.
1	32.864272	8.038391	1.571226	0.134253	1.282165	.
2	34.245280	9.073526	1.521404	0.329802	0.616704	.
Cluster	shutouts_pg	saves_pg	outs_pitches_pg	hits_allowed_pg	\	
0	0.063912	0.221949		26.815127		8.795287
1	0.006667	0.234904		25.836590		8.038659
2	0.061725	0.152523		26.722419		9.085837
Cluster	homeruns_allowed_pg	walks_allowed_pg	strikeouts_by_pitcher			
0		0.878630	3.218613			6.10
1		1.283065	3.392031			8.67
2		0.616873	3.024662			4.77
Cluster	errors_pg	double_plays_pg	fielding_percentage_pg			
0	0.806322		0.898966			0.006075
1	0.577261		0.797318			0.016430
2	1.316525		0.847623			0.007685

[3 rows x 27 columns]

Top features (highest variance between clusters):

games_played	2667.326984
strikeouts_by_pitchers_pg	3.931369
strikeouts_by_batters_pg	3.911271
at_bats_pg	0.548357
hits_allowed_pg	0.292248
outs_pitches_pg	0.291804
hits_pg	0.287299
errors_pg	0.143215
rank	0.142572
homeruns_allowed_pg	0.112649

dtype: float64

ANOVA Results (most significant differences):

	F-statistic	p-value
fielding_percentage_pg	18736.816379	0.000000e+00
games_played	8871.794288	0.000000e+00
outs_pitches_pg	181.691770	1.040376e-71

```
errors_pg           114.747947  3.061054e-47
strikeouts_by_batters_pg    111.933802  3.580371e-46
strikeouts_by_pitchers_pg   111.346058  5.989883e-46
triples_pg          83.404598  3.756219e-35
complete_games_pg    83.041942  5.213891e-35
```

```
df['PCA1'] = X_pca[:, 0]
df['PCA2'] = X_pca[:, 1]

# Export to CSV
output_file_path = "baseball_clusters.csv"
df.to_csv(output_file_path, index=False)

print(f"Data exported to {output_file_path}")
```

```
Data exported to baseball_clusters.csv
```